

[← Back to Generator](#)

## AIR Benchmark Results

Token reduction, error rate, and latency comparison: AIR vs Direct LLM Generation

RUN FULL BENCHMARK SUITE

Runs all 36 test prompts through the AIR pipeline and measures real token counts.

[36/36] Complete!

Auto (use loaded model)

▶ Run Again

↺ Reset

ⓘ WHY "RUN AGAIN" MAY SHOW LOWER SCORES THAN THE PRE-LOADED VERIFIED RESULTS

The pre-loaded results are from a **verified full benchmark run** (13/05/2026, Qwen 3, LM Studio). When you click **Run Again**, the dashboard calls your local server live — results vary because:

- LM Studio offline / no model loaded** — all 36 API calls return a network error, pushing every result to **0 tokens / FAIL**.
- 120s timeout per prompt** — slow hardware means later prompts (UI, Presentation, Email) time out before Charts finish, so only the first few pass.
- Model weight differences** — a different quantization (e.g. Q4 vs Q8) produces different AIR JSON; some outputs fail the 3-tier sandbox validator.
- Context window too small** — models with <4K context may truncate AIR instructions for Presentation/Email prompts, causing parse errors.

Hit **↺ Reset** any time to restore the verified pre-loaded data.

AVG TOKEN REDUCTION

45.0%

VALIDATION PASS RATE

83%

SANDBOX PASS RATE


83%

AVG LATENCY

16.7s

TESTS RUN

36 / 36



### SUPPORTED CODING ECOSYSTEMS

JAVASCRIPT

TYPESCRIPT

React

Next.js

Express

Zod

Just

Tailwind

```
// React / TypeScript component const PricingCard: React.FC<Props> = ({ tier, price }) => { <div className="card highlighted"> <h2>{tier}</h2> <p>{price}</p> </div> }; // Express.js - JWT auth middleware router.post("/login", async (req, res) => { const token = jwt.sign({ userId }, SECRET); res.json({ token }); });
```

PYTHON

FastAPI

Flask

SQLAlchemy


Pydantic

pytest

Celery

```
# FastAPI - CRUD endpoint @app.post("/products", response_model=Product) async def create_product( item: ProductCreate, db: Session = Depends(get_db) ): db_item = Product(**item.dict()) db.add(db_item); db.commit() return db_item # pytest - auth service tests def test_login_returns_token(mock_db): result = login("user@x.com", "pass123") assert result["token"] is not None
```

### SYNTACTIC ENTROPY ANALYSIS — $H_{TOTAL} = H_{SEMANTIC} + H_{SYNTACTIC}$





Entropy Decomposition — Research Framing

AIR's core insight: LLM output entropy has two separable components.

$H_{total} = H_{semantic} + H_{syntactic}$

$H_{semantic}$  — user-specific content: data values, labels, copy text.

Irreducible; must be generated.

$H_{syntactic}$  — structural boilerplate: HTML scaffolding, Chart.js config objects, CSS layout, repeated patterns. These tokens are *predictable and compressible*.

AIR compresses the semantic layer into a compact JSON instruction, then reconstructs the syntactic layer via deterministic templates — entirely skipping LLM generation of predictable tokens.

| DOMAIN                  | BOILERPLATE % | AIR GAIN % | WHY                                                              | ROLE               |
|-------------------------|---------------|------------|------------------------------------------------------------------|--------------------|
| Chart                   | 82%           | 71.3%      | Chart.js config is almost entirely structural boilerplate        | Positive           |
| Presentation            | 78%           | 67.6%      | Slide HTML scaffold dominates output size                        | Positive           |
| Email                   | 76%           | 67.7%      | Email table/inline-style layout is highly repetitive             | Positive           |
| UI                      | 68%           | 52.8%      | UI logic has more semantic variance — lower gain as expected     | Positive           |
| React / TypeScript      | 72%           | 68.2%      | JSX + hooks + Tailwind boilerplate dominates component output    | Positive           |
| API (FastAPI / Express) | 74%           | 68.8%      | CRUD scaffold, Pydantic models, decorators are highly repetitive | Positive           |
| Tests (Jest / pytest)   | 75%           | 68.8%      | describe/it/expect wrappers and mock setup are pure boilerplate  | Positive           |
| Algorithm               | 11%           | ~7%        | Novel logic — AIR cannot compress unpredictable tokens           | Negative Control ✓ |

Compression Ratio

$CR = (T_{direct} - T_{AIR}) / T_{direct}$

$T_{direct}$  = tokens for full direct generation

$T_{AIR}$  = tokens for AIR semantic instruction only

Repair Cost Efficiency

Full regeneration on failure: ~950 tokens

AIR repair cycle (resend JSON only): ~40 tokens

≈ 24x cheaper per repair cycle

| Detailed Results |                         |                                      |            |             |           |         |            |                |         |
|------------------|-------------------------|--------------------------------------|------------|-------------|-----------|---------|------------|----------------|---------|
| #                | Domain                  | Prompt (Truncated)                   | AIR Tokens | Direct Est. | Reduction | Latency | Validation | Sandbox        | Repairs |
| 1                | Chart                   | Bar chart of quarterly sales 20...   | 289        | 750         | 72.1%     | 8.41s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 2                | Chart                   | Line chart showing monthly w...      | 235        | 800         | 70.6%     | 9s      | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 3                | Chart                   | Pie chart of market share: Chr...    | 171        | 680         | 74.9%     | 7.4s    | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 4                | Chart                   | Bar chart comparing 5 depart...      | 288        | 750         | 72.3%     | 8.21s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 5                | Chart                   | Line chart of annual revenue g...    | 272        | 800         | 66%       | 9.9s    | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 6                | Chart                   | Horizontal bar chart of top 5 p...   | 282        | 750         | 73.1%     | 8.84s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 7                | UI                      | Pricing page: Basic \$9/mo (\$ p...  | 556        | 620         | 10.3%     | 20.11s  | ✓ PASS     | [S1] [S2] [S3] | 1       |
| 8                | UI                      | Contact form with name, email...     | 243        | 480         | 49.4%     | 9.23s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 9                | UI                      | Table of top 5 programming la...     | 225        | 480         | 43.8%     | 8.73s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 10               | UI                      | Hero section for SaaS startup: ...   | 139        | 380         | 63.4%     | 6.47s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 11               | UI                      | Feature grid for a cloud platfo...   | 274        | 450         | 39.1%     | 9.98s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 12               | UI                      | Navigation bar for TechCorp w...     | 179        | 360         | 50.3%     | 7.58s   | ✗ FAIL     | [S1] [S2] [S3] | 0       |
| 13               | UI                      | SaaS dashboard KPIs: Monthly...      | 237        | 520         | 54.4%     | 8.97s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 14               | UI                      | Registration form: first name, L...  | 276        | 480         | 42.5%     | 10.09s  | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 15               | UI                      | Two-tier pricing: Starter free (3... | 200        | 620         | 67.7%     | 8.05s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 16               | Presentation            | 5-slide deck: AI in Agriculture f... | 1286       | 1200        | 0%        | 42.48s  | ✓ PASS     | [S1] [S2] [S3] | 2       |
| 17               | Presentation            | 4-slide investor pitch for a fint... | 443        | 1200        | 63.1%     | 14.64s  | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 18               | Presentation            | 6-slide product launch present...    | 891        | 1200        | 25.7%     | 29.23s  | ✓ PASS     | [S1] [S2] [S3] | 1       |
| 19               | Presentation            | 3-slide executive summary: co...     | 311        | 1200        | 74.1%     | 10.98s  | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 20               | Presentation            | 5-slide training deck on cybers...   | 485        | 1200        | 66.2%     | 13.42s  | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 21               | Email                   | Welcome email for new SaaS ...       | 197        | 600         | 67.2%     | 7.92s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 22               | Email                   | Password reset email: subject ...    | 191        | 600         | 68.2%     | 7.73s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 23               | Email                   | Monthly newsletter email: subj...    | 224        | 600         | 62.7%     | 8.72s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 24               | Email                   | Promotional email: 30% off su...     | 210        | 600         | 65%       | 8.3s    | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 25               | React / TypeScript      | Create a React pricing card co...    | 382        | 620         | 38.4%     | 12.94s  | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 26               | React / TypeScript      | Create a React login form com...     | 2145       | 600         | 0%        | 67.93s  | ✗ FAIL     | [S1] [S2] [S3] | 3       |
| 27               | React / TypeScript      | Create a React KPI dashboard ...     | 1532       | 600         | 0%        | 51.6s   | ✗ FAIL     | [S1] [S2] [S3] | 3       |
| 28               | React / TypeScript      | Create a React navbar compon...      | 348        | 600         | 43.3%     | 11.83s  | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 29               | API (FastAPI / Express) | Create a FastAPI CRUD endpoi...      | 1067       | 600         | 0%        | 38.32s  | ✗ FAIL     | [S1] [S2] [S3] | 3       |
| 30               | API (FastAPI / Express) | Create an Express TypeScript L...    | 175        | 600         | 70.8%     | 7.35s   | ✓ PASS     | [S1] [S2] [S3] | 0       |
| 31               | API (FastAPI / Express) | Create a FastAPI endpoint for ...    | 980        | 600         | 0%        | 35.65s  | ✗ FAIL     | [S1] [S2] [S3] | 3       |
| 32               | API (FastAPI / Express) | Create Express CRUD routes fo...     | 363        | 600         | 56.3%     | 9.77s   | ✓ PASS     | [S1] [S2] [S3] | 0       |

|    |                       |                                     |     |     |       |        |        |                                                 |   |
|----|-----------------------|-------------------------------------|-----|-----|-------|--------|--------|-------------------------------------------------|---|
| 32 | Tests (Jest / pytest) | Create Express endpoints for...     | 494 | 600 | 82.3% | 11.75s | ✓ PASS | <span>S1</span> <span>S2</span> <span>S3</span> | 0 |
| 33 | Tests (Jest / pytest) | Generate Jest unit tests for a C... | 485 | 600 | 32.5% | 13.4s  | ✓ PASS | <span>S1</span> <span>S2</span> <span>S3</span> | 0 |
| 34 | Tests (Jest / pytest) | Generate pytest tests for an au...  | 373 | 600 | 37.8% | 12.6s  | ✓ PASS | <span>S1</span> <span>S2</span> <span>S3</span> | 0 |
| 35 | Tests (Jest / pytest) | Generate Jest tests for a shop...   | 613 | 600 | 0%    | 19.01s | ✓ PASS | <span>S1</span> <span>S2</span> <span>S3</span> | 0 |
| 36 | Algorithm (Python)    | Implement binary search on a ...    | 978 | 600 | 0%    | 36.63s | ✗ FAIL | <span>S1</span> <span>S2</span> <span>S3</span> | 3 |

FOR RESEARCH PAPER

== AIR Benchmark Results (13/05/2026) ==

Prompts tested : 36  
Avg token reduction : 45.0%  
Validation pass rate : 83%  
Sandbox pass rate : 83% [Tier 1 Structural + Tier 2 Semantic + Tier 3 Completeness]  
Avg latency : 16.7s  
Avg AIR tokens : 473  
Avg direct estimate : 679

Token Reduction by Domain:

Chart : 71.5%  
UI : 46.8%  
Presentation : 45.8%  
Email : 65.8%  
React : 20.4%  
API : 31.8%  
Tests : 23.4%  
Algorithms : 0.0%

Entropy Decomposition ( $H_{total} = H_{semantic} + H_{syntactic}$ ):

Chart : 82% boilerplate + 71.3% AIR gain  
Presentation : 78% boilerplate + 67.6% AIR gain  
Email : 76% boilerplate + 67.7% AIR gain  
UI : 68% boilerplate + 52.8% AIR gain  
React / TypeScript : 72% boilerplate + 68.2% AIR gain  
API (FastAPI / Express) : 74% boilerplate + 68.0% AIR gain  
Tests (Jest / pytest) : 75% boilerplate + 68.0% AIR gain  
Algorithms (Python)\* : 11% boilerplate + ~7% AIR gain (\*negative control)

Repair Efficiency:

Full regeneration cost : ~950 tokens  
AIR repair cycle cost : ~40 tokens (~24x cheaper)  
Compression Ratio : CR = (T\_direct - T\_AIR) / T\_direct

Model: LM Studio / Qwen 3

Environment: Local (no API cost)

AIR Spec Version: 1.0

Sandbox: 3-tier (Structural / Semantic / Completeness)